

Lab 5: Orchestrate Specialist Subagents

Business Innovation with AI Agents · TGS-2020503395 · 45 minutes

Objective

Split one campaign problem into isolated research, audience and risk tasks with structured evidence handoffs.

Deliverable

A task graph, three specialist handoffs, conflict-resolution note and human decision record.

Prerequisites

- A trainer-approved laptop and classroom workspace.
- Use only fictional data supplied in this lab.
- Do not paste API keys, customer data or personal information into evidence.
- Lab 4 skill and memory distinction.

Procedure

Step 1. Read campaign-brief.md and define three independent specialist outputs.

Step 2. Complete task-graph.md with parents, dependencies, acceptance criteria and stop conditions.

Step 3. Use `delegate_task` in Hermes or the trainer demonstration to run research, audience and risk specialists in isolated contexts.

```
hermes chat --toolsets file,web,delegation
```

Step 4. Require every specialist to return the handoff schema in handoff-template.md.

Step 5. Do not give specialists publication, messaging or spending tools.

Step 6. Compare handoffs. If claims conflict, keep both evidence trails and route them to the reviewer.

Step 7. Synthesize only after all parent tasks are complete; record missing evidence as a blocker.

Step 8. Complete human-decision.md and run the verifier.

```
python3 verify.py
```

Acceptance test

Three role-bounded handoffs exist, conflicts remain visible, the synthesis waits for dependencies and the human decision record controls release.

Evidence checklist

- ☐ Commands/configuration with secrets redacted
- ☐ Screenshot or file evidence
- ☐ Acceptance test result
- ☐ Failure and resolution
- ☐ Reflection and improvement

Safety and cleanup

- Use fictional data only.
- Never submit .env, tokens, keys or personal information.
- Stop classroom gateways and remove recurring jobs after evidence capture.